

VUG7 Keynote:

Dude! Where's my Varnish 4.0 ?



## Distractions & Challenges

Poul-Henning Kamp

[phk@Varnish.org](mailto:phk@Varnish.org)

@bsdphk

Things my great grandmother said:

”Kanføre og gæstebud ta’r man som det kommer”

Roughly translated:

Sleigh-weather and parties happen when they do











TIL SALG

2 store grundstykker

2.776 og 6.569 m<sup>2</sup>

Til boligformål

hørme ERIKVERV

Tlf. 58 58 08 40




















# Flag days vs. regular days

Things we can do in major releases:

- Change how Varnish works

- Change VCL syntax

- Change VMOD APIs

- Change important defaults

Things we can do in minor releases:

- Fix bugs

- Add VMODs

- Speed things up



## Varnish 4: Changing assumptions

HTTP/1.0: Serialized communication

req — resp — req — resp — res — resp

HTTP/2.0: Multiplexed communication

req1+req2+req3 — resp2 — req4 — resp1,resp3,push1



# Varnish model for HTTP/1.0:

```
recv req
|
miss ? -> fetch
|           |
deliver <---+
|
done
```



# Varnish model for HTTP/2.0:

?

HTTP/2.0 doesn't exist yet

And may not exist for some years

And may not be any good once it exists



# Varnish model for HTTP/2.0:

"session thread" vs. "state machine"  
- depends on protocol design.

Per request thread  
- Runs the VCL-obstacle course  
- Triggers fetches

Per fetch thread  
- Background / parallel / streaming



Req thread

Fetch thread

Miss:

Allocate busyobj+objcore  
vcl\_miss{}

-----> Build bereq

<----- release req

Deliver grace/  
or

vcl\_backend\_fetch{}

send request

rx response (hdrs)

vcl\_backend\_response{}

Wait for valid obj <-----

mark "valid"

Streaming delivery <-----

rx body

\*done\*



Req thread

Fetch thread

Pass:

Allocate busyobj+objcore  
vcl\_pass{}

-----> Build bereq  
vcl\_backend\_fetch{  
send request  
send req.body

. <----- release req  
. rx response (hdrs)  
. vcl\_backend\_response{}

Wait for valid obj <----- mark "valid"

Streaming delivery <----- rx body  
\*done\*

```
sub vcl_backend_fetch {  
    return (fetch);  
}
```



```
sub vcl_backend_response {  
    if (beresp.ttl <= 0s ||  
        beresp.http.Set-Cookie ||  
        beresp.http.Surrogate-control ~ "no-store" |  
        (!beresp.http.Surrogate-Control &&  
            beresp.http.Cache-Control ~  
                "no-cache|no-store|private")) ||  
        beresp.http.Vary == "*") {  
        /*  
        * Mark as "Hit-For-Pass"  
        * for the next 2 minutes  
        */  
        set beresp.ttl = 120 s;  
        set beresp.do_pass = true;  
    }  
    return (deliver);  
}
```

```
sub vcl_lookup {  
    if (obj.uncacheable) {  
        return (pass);  
    }  
    if (obj.ttl >= 0s) {  
        return (deliver);  
    }  
    if (obj.ttl + obj.grace > 0s) {  
        return (deliver_stale);  
    }  
    return (deliver);  
}
```



# VMOD objects

```
sub vcl_init {  
    new foo = directors.random();  
    foo.add_backend(s1, 1);  
    foo.add_backend(s2, 1);  
}  
  
sub vcl_recv {  
    set req.backend = foo.backend();  
}
```

---

```
Object random() {  
    Method VOID .add_backend(BACKEND, REAL)  
    Method BACKEND .backend()  
}
```

# VCL context structure

We used to use "struct req\*" as the root pointer

(We'll always have a request, right ?)

With vcl\_backend\_(fetch|response){} we have a struct "busyobj\*" as root pointer.

So now we need to pass wrk\*,req\*,busyobj\* and ...

Solve this with a new "struct vrt\_ctx" container



```

struct vrt_ctx {
    unsigned
#define VRT_CTX_MAGIC
    unsigned
    unsigned
    struct vs1_log
    struct VCL_conf
    struct ws
    struct req
    struct http
    struct http
    struct http
    struct busyobj
    struct http
    struct http
    magic;
    0x6bb8f0db
    method;
    *handling;
    *vs1;
    *vcl;
    *ws;
    *req;
    *http_req;
    *http_obj;
    *http_resp;
    *bo;
    *http_bereq;
    *http_beresp;
};

```

# Varnish log filtering

```
param.show vsl_mask
```

```
200 373
```

```
vsl_mask      -VCL_trace,-WorkThread,-Hash
```

Default is default

Mask individual VSL messages from being logged. default Set default value

Use +/- prefixe in front of VSL tag name, to mask/unmask individual VSL messages.

```
param.set vsl_mask +vcl_trace,-expkill
```

```
200 0
```

```
param.show vsl_mask
```

```
200 371
```

```
vsl_mask      -ExpKill,-WorkThread,-Hash
```

```
[...]
```

# Workspace changes

Used to be:

- Requests have private workspace

- Allocate all dynamic space on thread-stack

- Including all space for backend-fetch

=> Even 99.99% hit rate will cause all threads to drag around space for backend fetches.



# Workspace changes

Changed to:

Requests have private workspace  
Backend requests ("busyobj") ditto  
Thread stack is used minimally

Why:

$N_{\text{threads}} > N_{\text{req}} \gg N_{\text{busyobj}}$   
= More threads use less memory

# Purge been properly adopted

```
sub vcl_recv {  
    if (req.method == "PURGE") {  
        return (purge);  
    }  
}
```

```
sub vcl_purge {  
    /* XXX: Could decide to fetch here */  
    return (error(200, "Purged"));  
}
```

# The new varnishlog/ncsa stuff

Overall idea: Tcpdump like filtering language

Also: Track the req/fetch changes in varnishd

Ask Martin, not me :-)



# Varnish penetration ?

varnish site:netcraft.com

Internet Billeder Maps Indkø

Ca. 43.000 resultater (0,08 sekunder)

503 guru meditation xid


Internet Billeder Maps Ind

Ca. 62.100 resultater (0,12 sekunder)

Interesse over tid ?

Tallet 100 repræsenterer den største søgeinteresse

☒ Nyhedsoverskrifter ☐ Prognose ?



[http://trends.builtwith.com/:](http://trends.builtwith.com/)

## Websites using Varnish


Get a list of websites using Varnish. We know of **471,877 websites using Varnish**. 99,030 websites within the most visited sites on the internet and an additional **372,847** websites on the rest of the web.

## Top sites using Varnish

Domain	SIRank
<a href="#">pinterest.com</a>	
<a href="#">Full Profile</a>	
<a href="#">answers.com</a>	
<a href="#">Full Profile</a>	
<a href="#">weather.com</a>	
<a href="#">Full Profile</a>	
<a href="#">go.com</a>	
<a href="#">Full Profile</a>	
<a href="#">examiner.com</a>	
<a href="#">Full Profile</a>	
<a href="#">nydailynews.com</a>	
<a href="#">Full Profile</a>	
<a href="#">merriam-webster.com</a>	
<a href="#">Full Profile</a>	
<a href="#">grindtv.com</a>	
<a href="#">Full Profile</a>	
<a href="#">people.com</a>	
<a href="#">Full Profile</a>	
<a href="#">usmagazine.com</a>	
<a href="#">Full Profile</a>	

## Varnish Usage Trends

Varnish is a web accelerator / reverse proxy caching server.



# VML: A four step programme

VML = Varnish Moral License

- 1 Happy users ask for VML
- 2 I send them an invoice
- 3 They pay the invoice
- 4 I develop more Varnish

# VML Status so far:

2010: 900 hours  
2011: 924 hours  
2012: 586 hours  
2013: 333 hours (jan-may)

Thank you!

Also host servers  
Builds releases  
Organizes VUGs

- \* Varnish software
- \* GLOB0 Communications
- \* UPLEX
- \* (You company could be here!)



VML going forward

Real-estate solves tax-limit (<33%)

Looking to ramp up Varnish share of work

Looking for more licenses:

room for 6000-8000 EUR|USD pr. month

One-time/recurring: your choice

<http://phk.freebsd.dk/VML/>