


How I spent my summer vacation


Poul-Henning Kamp

phk@FreeBSD.org

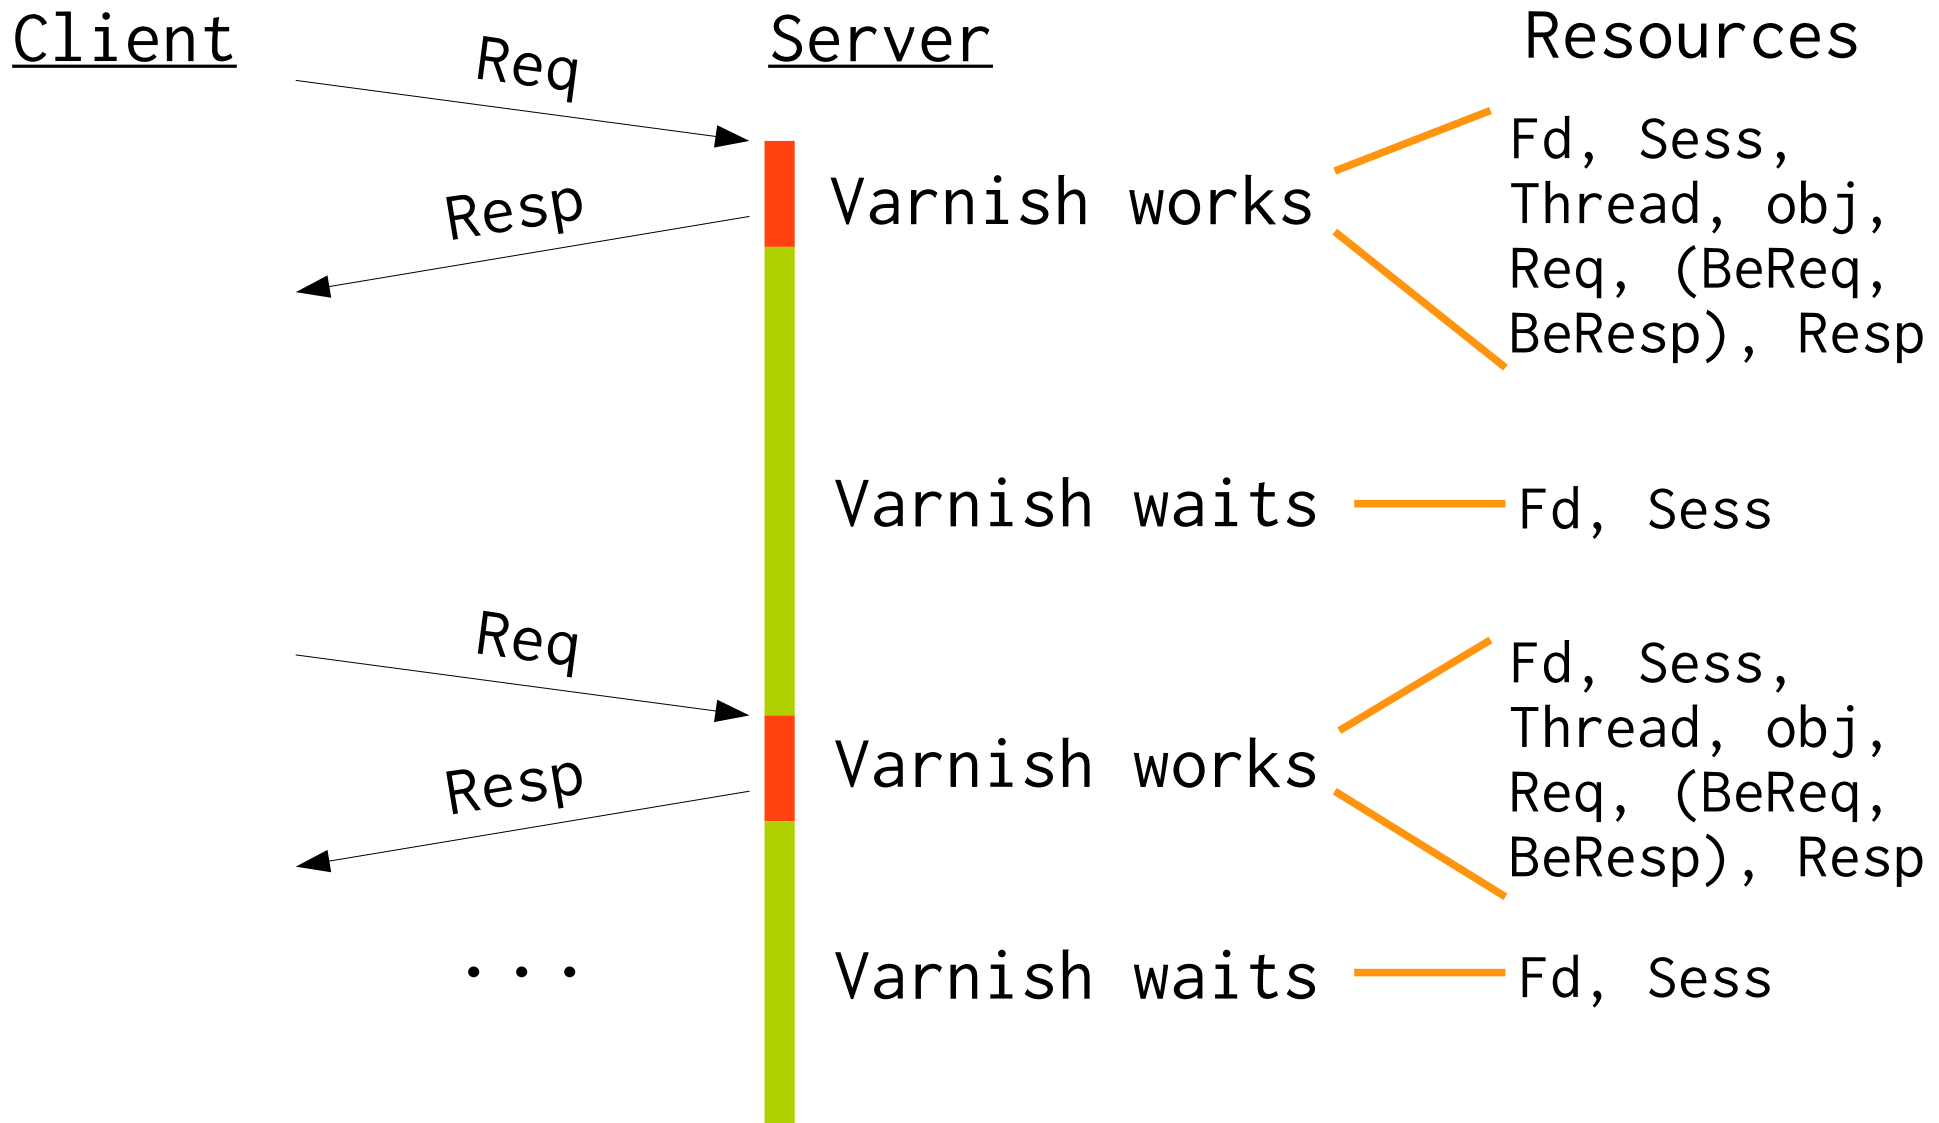
”HTTP is pure request->response”




”HTTP is pure request->response”



”HTTP is pure request->response”




Work



Time


”High aspect ratio, farm friendly job”

Work



Time

”High aspect ratio, farm friendly job”



”On average, Varnish doesn’t do s**t”
— sky

HTTP design scenario

Open TCP

GET dilbert.com::index.html

200 <some bytes>

Close TCP

<Render page>

Open TCP

GET dilbert.com::dilbert.19940622.gif

200 <some bytes>

Close TCP

<ReRender page>

HTTP today's scenario

Open TCP

GET dilbert.com::index.html

200 <some bytes>

GET dilbert.com::favicon.ico

200 <some bytes>

GET dilbert.com::bg_body.jpg

200 <some bytes>

GET dilbert.com::ico_new.png

200 <some bytes>

GET dilbert.com::bg_strip_header_featured.gif

200 <some bytes>

[42 other objects]

HTTP today's scenario

Open TCP

GET dilbert.com::index.html

200 <some bytes>

GET dilbert.com::favicon.ico

200 <some bytes>

GET dilbert.com::bg_body.jpg

200 <some bytes>

GET dilbert.com::ico_new.png

200 <some bytes>

GET dilbert.com::bg_strip_headire

200 <some bytes>

[42 other objects]

Cookies

$\frac{1}{2}$ RTT's

2

1

1

1

1

1

1

1

1

1

1

1

1

1

1

1

1

1

1

42

2×42

HTTP creativity boundary

Page render time is dominated by speed of light

→ 299,792,458 m/s: It's the law.

Supposedly unchangeable invariants:

- TCP/IP
- TCP port 80 (& 443)
- HTML/CSS/JS model

HTTP/2.0 pipeline/mux

Open TCP

GET dilbert.com::index.html

200 <some bytes>

GET dilbert.com::favicon.ico

GET dilbert.com::bg_body.jpg

GET dilbert.com::ico_new.png

GET dilbert.com::bg_strip_header_featured.png

[42 more gets]

220 favicon.ico <some bytes>

220 ico_new.png <some bytes>

220 bg_body.jpg <some bytes>

220 bg_strip_header_featured.png <some bytes>

[42 more 220's]

Cookies

$\frac{1}{2}$ RTT's

1

2

1

1

1

HTTP/2.0 server push

Open TCP

GET dilbert.com::index.html

200 <some bytes>

220 favicon.ico <some bytes>

220 ico_new.png <some bytes>

220 bg_body.jpg <some bytes>

220 bg_strip_header_featured.g <some bytes>

[42 other objects]

Cookies

$\frac{1}{2}$ RTT's

1

2

1

HTTP/2.0 political wishlist

Mandatory SSL/TLS

- + End-to-end privacy
- Kills all caches
- Law mandates inspection (jails, kids ...)
- Killed, but a recurring zombie issue

Session-IDs

- Trying to avoid cookies

Proxy HTTPS handling

- How to do inspection on HTTPS

HTTP/2.0 technical wishlist

Compression

→ Really about cookies

Better protocol serialization

→ Faster, safer processing

UniCode

→ Really about avoiding encodings

HTTP/2.0 appearantly not about

HTTP → HTTPS upgrade cost

→ Connection upgrade would have been nice

Proxy mux efficiency

→ Mux secured and unsecured traffic

HTTP routers

→ Routable easily accessible envelope

Handling realistic future bandwidths

→ 10 years → 10Gb/s on laptops

→ 20 years → 100Gb/s on laptops

HTTP/2.0 — Because IPv6 was such a success

[RANT]

Shooting fish in a waterfall

HTTP/2.0 is an uncertain target

- No realistic timeline
- Very unclear featureset
- May not even catch on when done

What does SPDY do in the meantime ?

What should Varnish do ?

Wouldn't it be nice...

What *could* Varnish do with/for HTTP/2.0 ?

- Perfect for protocol migration
1.0 backend, 2.0 client or vice versa
- Add server-push to HTTP/1.0 backend
VCL good for expressing policy & logic

```
sub vcl_recv {  
    if (req.proto == "HTTP/2.0" && req.url == "/") {  
        push("style.css");  
        push("favicon.ico");  
        push("fp.js");  
    }  
}
```

Where are we ?

Varnish ≤ 3 heavily optimized for HTTP/1.1

Technical

- Threads, pools, workspaces

Architectural

- VCL request flow
- VSL log record contents

Where do we want to end up ?

Multiple concurrent requests per session

Multiprotocol: HTTP/1.1, SPDY, HTTP/2.0 ...

Decouple client & backend transactions

Separate thread does backend transaction

Introduce rendez-vous between threads

- refcount/copy req.* ?
- Original session may be gone before beresp.*
- How does this look in VCL ?

Not a lot of work, but some pitfalls & tarpits

Decouple client & backend transactions

Separate thread does backend transaction

HTTP/1.1 Payoff:

- + cleaner streaming
- + parallel ESI fetches
- + no-delay grace mode

Multiprotocol

The protocol is what defines a session

→ VCL not protocol agnostic
ie: setting "Connection: Close" in vcl_pipe{}
What does "pipe" mean for HTTP/2.0 anyway ?

→ Per protocol VCL ?

→ Or VCL just does "abstract" request ?

Where does session trickery go then ?

Multiprotocol

The protocol is what defines a session

HTTP/1.1 session/protocol code:

- accept connections (VCA)
- receive & validate requests (HTTP)
- feed in requests to be handled (HTTP1_FSM)
- responses in abstract form (RSP)
- serializes response onto connection (WRW)

... All this must be objectified
= Design OO-API

Multiprotocol

The protocol is what defines a session

HTTP/2.0 session/protocol code:

- accept connections (VCA)
Reuse mandatory (due to UPGRADE)
- receive & validate requests (HTTP2)
May be best with per-session rx-thread
- feed in requests to be handled (HTTP2_FSM)
- responses in abstract form (RSP)
- serializes response onto connection (WRW)
Probably better with a per-session thread
rather than per-object thread.

Multiprotocol

The protocol is what defines a session

HTTP/1.1 Payoff:

... none really

Lots of work

But show stopper for multiprotocol support

... We might as well start early

VCL language design is the hard part

VCL is a "domain specific language"

→ Strong expression of domain mechanics

→ Hide/Handle invariant details

→ Minimize boilerplate code

VCL language design is the hard part

VCL is also where users invest resources

→ Changes not welcome, unless they are very good

But VCL is why Varnish is popular

→ Must stay relevant and on top of game

VCL language vs. magic buttons

We are trending towards magic buttons:

do_gzip, do_gunzip, do_esl, always_miss

N^2 complexity of interaction

→ Hard to define

→ Hard for users to remember

Prefer programmatic model in VCL

→ Unfortunately leads to more VCL code

PURGE with magic buttons

```
sub vcl_recv {  
    if (req.request == "PURGE" && client.ip ~ pac1) {  
        set req.do_purge_all_variants;  
    }  
}
```

PURGE with semi-programmatic model:

```
sub vcl_recv {  
    if (req.request == "PURGE && client.ip ~ pac1) {  
        return (lookup);  
    }  
}  
sub vcl_lookup {  
    if (req.request == "PURGE) {  
        purge_all_variants;  
        return(synth(200, "Purge done"));  
    }  
}
```


PURGE with full programmatic model:

```
sub vcl_lookup {  
    if (req.request == "PURGE") {  
        for_variants {  
            if (backend.default.healthy) {  
                purge(obj);  
            } else {  
                set obj.ttl = 0s;  
                set obj.grace = 10m;  
            }  
        }  
        return(synth(200, "Purge done"));  
    }  
}
```

And yeah, I'm not asking you to decide...

"One of each, please ?"

Finally:

A big thanks to the Varnish Moral Licensees:

Varnish Software

Globo

UPLEX

FaceBook

They pay me for working on Varnish