



# Logging in Varnish 4



Martin Blix Grydeland  
Varnish Software

# Logging in Varnish

- Shared memory
- Always log everything (no debug switches)
- Utilities to inspect the log

# Shared memory log layout



# varnishlog in 3.0

- Groups log records into transactions
- -i and -x controls output
- -m <tag:regex> allows querying

# Problems with varnishlog in 3.0

- Only regex matching
- Only AND operator
- -i/-x and -m don't work together
- No req/bereq relationship
- Always copying
- Varnish 4 doesn't log by fd

# Logging goals for Varnish 4

- Transactions and transaction groups
- Query language
- Output controls
- Zero copy in the common case

# Varnish transactions

- A transaction is one work item in Varnish
- Share a single VXID
- Types of transactions
  - Client request
  - Backend request
  - ESI sub-request
  - Session

# Transaction groups ...

- ... group together transactions
- Grouping modes
  - VXID
  - Request
  - Session
  - Raw
- Names subject to change?



# Transaction groups are hierarchical

## Cache miss with request grouping

- Request
  - Backend request


## ESI included miss with session grouping

- Session
  - Request
    - ESI sub-request
      - Backend request
    - ESI sub-request
    - Backend request
  - Request
    - Backend request

# Cache miss with request grouping

- Request

○ Backend request



```
martin@friday:~$ varnishlog -g request
* << Request >> 32770
- Begin req 32769
- ReqMethod GET
- ReqURL /
- ReqProtocol HTTP/1.1
- ReqHeader TE: deflate,gzip;q=0.3 ...
- Link bereq 32771
- VCL_call DELIVER
- VCL_return deliver
- RespProtocol HTTP/1.1
- RespStatus 200
- RespResponse OK
- ReqEnd 1385330985.979025126
1385330985.978960991 -0.001315594 0.001251459
-0.001315594
- End
** << BeReq >> 32771
-- Begin bereq 32770
-- VCL_call BACKEND_FETCH
-- VCL_return fetch
-- BackendOpen 18 default(127.0.0.1,::1,8020)
127.0.0.1 45989
-- Backend 18 default default(127.0.0.1,::1,8020)
-- BereqMethod GET
-- BereqURL /
-- BereqEnd 1385330985.979187250
1385330985.980367422 0.000082792 0.000496101 0.000326045
0.000822146
-- End
```

# varnishlog invocation

```
varnishlog -g <grouping-mode> -q <query-expression>
```

Grouping mode is one of:

- vxid (default)
- request
- session
- raw

# Query language

- Operates on transaction groups
- Expr true if it matches one or more records
- Expr false if it matches no records

```
$ varnishlog -q 'respstatus < 500'
```

# Matching on transaction groups

```
$ varnishlog -g request \  
-q 'ReqURL eq "/"'
```

```
$ varnishlog -g request \  
-q 'Backend ~ default'
```

```
* << Request >> 32770  
- Begin req 32769  
- ReqMethod GET  
- ReqURL /  
- ReqProtocol HTTP/1.1  
- ReqHeader TE: deflate,gzip;q=0.3 ...  
- Link bereq 32771  
- VCL_call DELIVER  
- VCL_return deliver  
- RespProtocol HTTP/1.1  
- RespStatus 200  
- RespResponse OK  
- ReqEnd 1385330985.979025126  
1385330985.978960991 -0.001315594 0.001251459  
-0.001315594  
- End  
** << BeReq >> 32771  
-- Begin bereq 32770  
-- VCL_call BACKEND_FETCH  
-- VCL_return fetch  
-- BackendOpen 18 default(127.0.0.1,::1,8020)  
127.0.0.1 45989  
-- Backend 18 default default(127.0.0.1,::1,8020)  
-- BereqMethod GET  
-- BereqURL /  
-- BereqEnd 1385330985.979187250  
1385330985.980367422 0.000082792 0.000496101 0.000326045  
0.000822146  
-- End
```

# Usual suspects of operators

- String matching

`RespProtocol eq "HTTP/1.1"`

- Regular expressions

`ReqMethod ~ "GET|POST"`

- Integer matching

`RespStatus == 200`

`RespStatus >= 200`

- Float matching

`RespStatus == 200.`

# Boolean operators

- **and**

`RespStatus >= 500 and RespStatus < 600`

- **or**

`ReqURL == "/" or ReqURL == "/index.html"`

- **Group terms using parenthesis**

- **Negate using `not`**

# Other useful tricks

- Response time exceeds  $\frac{1}{2}$  second

```
ReqEnd[5] >= 0.5
```

- Client requests connection closed

```
ReqHeader:connection ~ close
```

- ESI miss (-g request)

```
{3+}Begin ~ Bereq
```



# Output controls

- Applied last, doesn't affect queries
- -i <taglist> / -I <taglist:regex>
- -x <taglist> / -X <taglist:regex>
- Taglists supports globbing (e.g. Req\*)

# Zero copy

- Everything done on shared memory
- Except when needed
- Transparently make copies
- Copy only what's needed
- Max time / number knobs



# varnishncsa in 4

- Query language
- Zero copy
- Format string parsed once
- Built for speed
- Feature parity with 3.0
  - Backend traffic logging planned

# Status of the other utilities

- varnishtop
- varnishhist
- varnishreplay

Not yet reimplemented using the new API

# API

- All of the magic happens in the API
- Better varnishd exceptional handling
- Easy integration

# API cont.

```
/* See varnish-cache/include/vapi/vsl.h for full documentation */
```

```
/* Create new query */
```

```
struct VSLQ *VSLQ_New(struct VSL_data *vsl, struct VSL_cursor **cp,  
    enum VSL_grouping_e grouping, const char *query);
```

```
/* API processing */
```

```
int VSLQ_Dispatch(struct VSLQ *vslq, VSLQ_dispatch_f *func, void *priv);
```

```
/* Callback function definition */
```

```
typedef int VSLQ_dispatch_f(struct VSL_data *vsl,  
    struct VSL_transaction * const trans[], void *priv);
```

# API cont.

```
struct VSL_transaction {
    unsigned                level;
    int32_t                 vxid;
    int32_t                 vxid_parent;
    enum VSL_transaction_e  type;
    struct VSL_cursor       *c;
};

struct VSL_cursor {
    struct VSLC_ptr         rec;
    const void              *priv_tbl;
    void                    *priv_data;
};

/* Traverse log entries */
int VSL_Next(struct VSL_cursor *c);
```

**Thanks for listening**