

Cloudy-cloud-enabled

Dynamic backends in Varnish 4.1



About me

- Developer
 - I contribute features, documentation and mostly bugs
 - I hate the web
 - I love Varnish
- French
 - ~~Your problem, not mine~~



More about me

- Dec 2011: I discover Varnish 3, I stop doing Java
- Jun 2012: I start my first VMOD
- Dec 2012: I give my first Varnish training
- Apr 2014: My first patch lands in Varnish
- Mar 2015: I join Varnish Software
- Oct 2015: First blog post at Varnish Software
- Dec 2015: My first VUG \o/

Two* things I learned

- There is almost always a caching requirement behind each functionality
- Simplicity is almost always preferred given more than one obvious solution
- Knowing the history of the project helps better understand its current state

The VBE and VDI subsystems



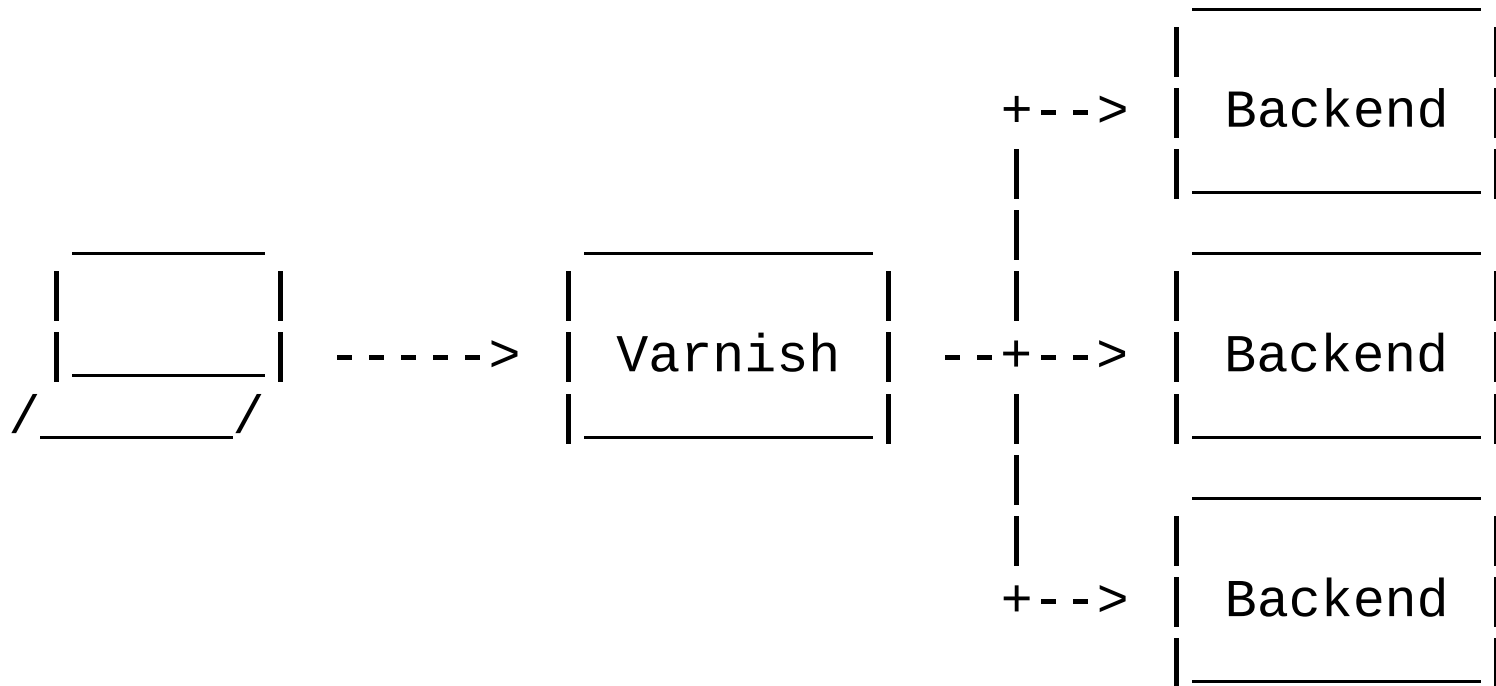
The VBE and VDI subsystems

- Back to basics



The VBE and VDI subsystems

- Back to basics



The VBE and VDI subsystems

- In Varnish 3, tight coupling with VCL

```
backend www {  
    .host = "www.example.com";  
    .port = "http";  
}
```

```
director dir random {  
    .retries = 5;  
    {  
        // We can refer to named  
        // backends  
        .backend = fs1;  
        .weight  = 7;  
    }  
    {  
        // Or define them inline  
        .backend = {  
            .host  = "fs2";  
        }  
        .weight  = 3;  
    }  
}
```


The VBE and VDI subsystems


- In Varnish 4, directors removed from VCL

```
backend www {  
    .host = "www.example.com";  
    .port = "http";  
}
```

```
import directors;  
  
sub vcl_init {  
    new dir =  
        directors.round_robin();  
    dir.add_backend(b1);  
    dir.add_backend(b2);  
}  
  
sub vcl_recv {  
    set req.backend_hint =  
        dir.backend();  
}
```

The VBE and VDI subsystems

- New capabilities
- But no more DNS director



The VBE and VDI subsystems

- Cloudy-cloud-enabled
 - VCL live-reload is immediate
 - Backends are static, VCL is dynamic
- How to cloud?
 - Listen to infrastructure events
 - Generate VCL
 - Reload VCL
 - Profit

Current state in 4.1



Current state in 4.1

- Changes in the director API
 - Native backends sort-of removed from VCL
 - VCL consumes the backend API
 - Same API available to VMODs
 - They can be dynamic as well as static
 - “Cluster” directors still available
 - Backend resolution chain handled by the VDI
 - “Custom” backends
 - No need to speak HTTP/1 over TCP on the backend side
 - Intermediary web servers can virtually disappear

Current state in 4.1

- Changes in the director API

```
/* Varnish 4.0 */
```

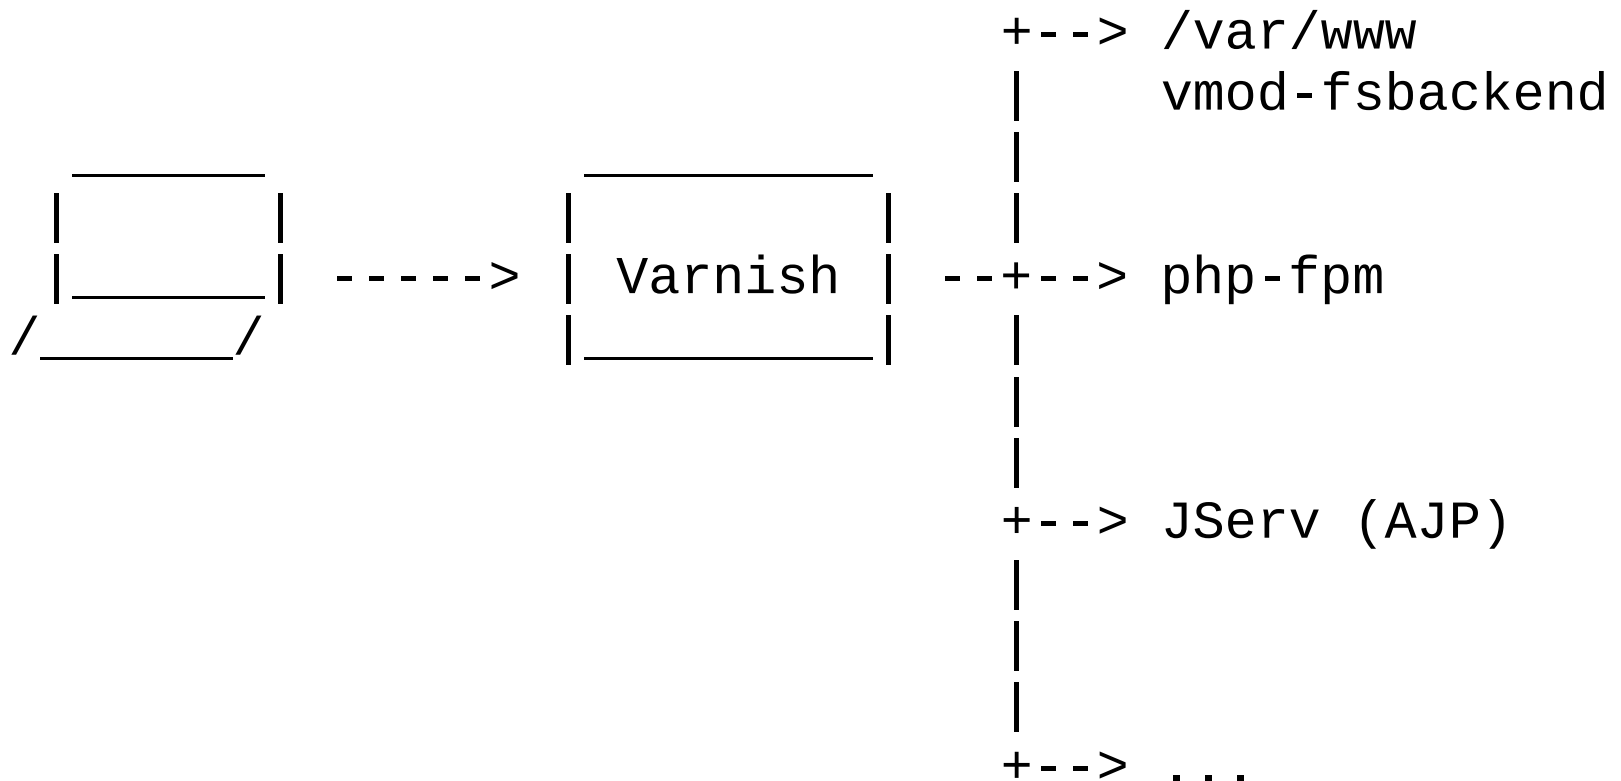
```
struct director {  
    unsigned    magic;  
    const char  *name;  
    char        *vcl_name;  
    vdi_getfd_f *getfd;  
    vdi_healthy *healthy;  
    void        *priv;  
};
```

```
/* Varnish 4.1 */
```

```
struct director {  
    unsigned    magic;  
    const char  *name;  
    char        *vcl_name;  
    vdi_healthy *healthy;  
    vdi_resolve_f *resolve;  
    vdi_http1pipe_f *http1pipe;  
    vdi_gethdrs_f *gethdrs;  
    vdi_getbody_f *getbody;  
    vdi_getip_f *getip;  
    vdi_finish_f *finish;  
    vdi_panic_f *panic;  
    void        *priv;  
    const void  *priv2;  
};
```

Current state in 4.1

- Custom backends in a nutshell
 - If you write a VMOD, you can:



Current state in 4.1

- Back to “VBE” backends
 - New VRT functions
 - VRT_new_backend
 - VRT_delete_backend
 - Create and delete them at any time*
 - No need to reload VCL
 - Don't delete backends you don't own
 - VCL backends are statically referenced for instance

DNS-based backends



DNS-based backends

- Problem to solve
 - Reload backends without reloading VCL
 - Infrastructure and cache policy have different life cycles
- Proposed solution
 - Rely on DNS to get a list of backends
 - ???
 - Profit

DNS-based backends

- Doesn't work with “VBE” backends

VCL may not compile

```
backend google {  
    .host = "google.com";  
    .port = "http";  
}
```

DNS-based backends

- Doesn't work with “VBE” backends
- A host name can resolve at most
 - One IPv4 address
 - One IPv6 address
- Varnish assumes both point to the same server

DNS-based backends

- Varnish 3
 - Create static backends
 - Enable/disable them according to look-ups

```
director directorname dns {  
    .list = {  
        .host_header = "www.example.com";  
        .port = "80";  
        .connect_timeout = 0.4s;  
        "192.168.15.0"/24;  
        "192.168.16.128"/25;  
    }  
    .ttl = 5m;  
    .suffix = "internal.example.net";  
}
```

DNS-based backends

- Varnish 4
 - Nothing
 - Might be possible though

DNS-based backends

- Varnish 4.1
 - <https://github.com/dridi/libvmmod-named>
 - Turns Varnish into a forward proxy...

```
probe healthcheck { ... }
```

```
sub vcl_init {  
    new dir = named.director("http");  
    dir.probe_with(healthcheck);  
    dir.set_ttl(5m);  
}
```

```
sub vcl_recv {  
    set req.backend_hint =  
        dir.backend("www.example.com");  
}
```

DNS-based backends

- Design differences
 - Static vs dynamic backends
 - Probe support in vmod-named
 - “White-list” support in the DNS director*
 - Not easily feasible in Varnish 4.1.0
 - Could be done by reusing ACLs
 - But my ACL patch missed the release window
 - Look-ups block workers in the DNS director
 - Look-ups are done in a separate thread in vmod-named
 - Production-ready vs PoC

Cloudy-cloud-enabled



Cloudy-cloud-enabled

- Varnish can now be aware of backends changes
- No need to reload the VCL for that
- Just sit back and relax
- Right?

Cloudy-cloud-enabled

- New capabilities, new caveats
- VCLs own backends instead of sharing
 - Because of dynamic backends*
 - Backends share connection pools
- A discarded VCL can stick for a long time*

Cloudy-cloud-enabled

- Let's play “spot the differences”

```
varnishstat -1 -f VBE.\* |  
awk '/vcls|happy/ {print $1 "\t" $2}'
```

```
/* Varnish 4.0 */
```

```
VBE.default(,::1,80).vcls 10  
VBE.default(,::1,80).happy 0
```

```
/* Varnish 4.1 */
```

```
VBE.vcl11.default.happy 0  
VBE.vcl12.default.happy 0  
VBE.vcl13.default.happy 0  
VBE.vcl14.default.happy 0  
VBE.vcl15.default.happy 0  
VBE.vcl16.default.happy 0  
VBE.vcl17.default.happy 0  
VBE.vcl18.default.happy 0  
VBE.vcl19.default.happy 0  
VBE.vcl110.default.happy 0
```

VCL Temperature



VCL Temperature

- VCL can be cold or warm
 - A VCL must be warmed up before use
 - It can eventually cool down after use
 - New knobs to play with
- Overall a breaking change
 - Plan your upgrade first
 - I mean it!

VCL Temperature

- A cold VCL
 - Should not get in the way of the active VCL
 - Should have the lowest possible footprint
 - Will drop native backends probes and stats
- Remember that
 - Backends will grow “linearly”
 - The more the probes, the more the overhead
 - The stats segment is limited

VCL Temperature

- Design is however not complete
 - VMODs may get the information too late
 - But creating a “cold” backend is not allowed
 - VCLs may warm up before completely cooling down
 - Known problem, a patch is ready*
- Currently
 - Dynamic backends work
 - But only in sub vcl_init{}
 - A workaround is to only have warm VCLs

Cloudy-cloud-enabled



Cloudy-cloud-enabled

- Better suited for backends in the cloud
 - VCL reloads have more constraints
 - But backends can be refreshed any time
 - Decouple infrastructure and policy
 - Integrate with your favorite stack
 - Expect more from future releases!
- A lot more powers to VMOD writers

Thank you for your time!
Questions?

Check out 4.1 VMODs on github:
- mbgrydeland/libvmod-fsbackend
- dridi/libvmod-named