

IMS

Backend Conditional Requests aka Cache Refresh/Refresh

VUG5 March 23, 2012

Geoffrey Simmons Nils Goroll





Agenda

- brief history
- experimental-ims
 - How it works
 - Refresh and storage
- Configuration
 - keep, stale_obj
 - vcl_stale(), vcl_refresh()?
- Too complex?

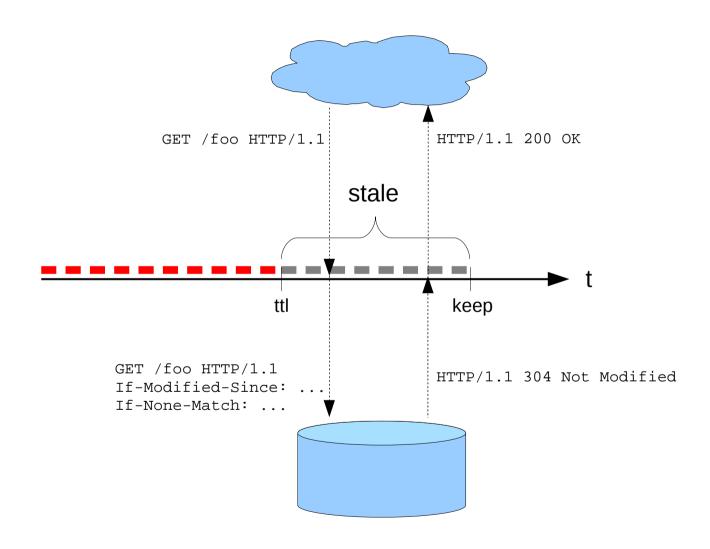
Brief History



- 2010-09-10
 - Received first Rackspace Implementation
- 2010-09-23
 - Proposed design on varnish-dev
- 2011-03-01
 - First patch posted
- Since 2011-06-02
 - Maintaining git branch experimental-ims

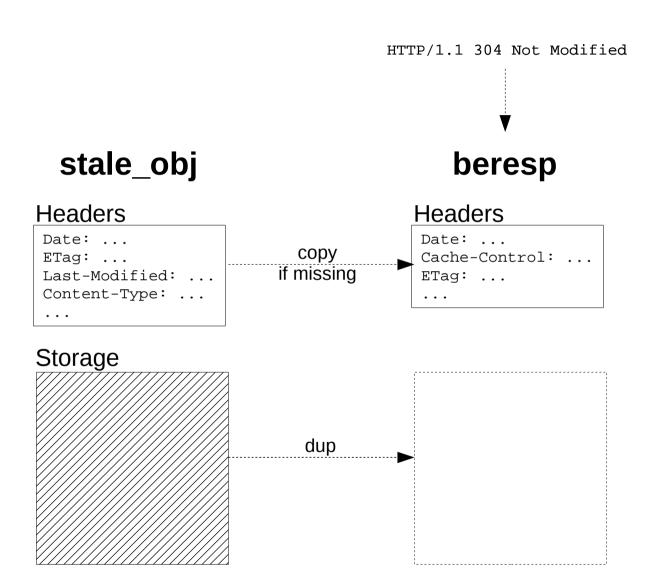


experimental-ims



Refresh & storage





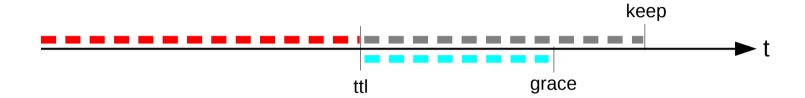
Storage dup



- dup == copy
 is unavoidable for -spersistent
- dup for other stevedores
 - Should be:
 - New object points to stale_obj's storage
 - Refcounts for storage
 - Currently in experimental-ims:
 - Copy for <u>all</u> stevedores



keep & grace



In experimental-ims:

- Both intervals begin after ttl
 - keep == grace, keep < grace, keep > grace all possible
- Expiration after ttl + max(keep, grace)
- grace has priority in overlapping intervals

stale_obj



Currently in experimental-ims:

```
sub vcl_miss {
    if (stale_obj) {
        set bereq.http.X-Foo = stale_obj.http.X-Bar;
    }
}
```

- All fields read-only
 - Otherwise just like obj
- In vcl_miss(), vcl_fetch(), vcl_error()
 - May or may not exist!

stale_obj



```
# DON'T DO THIS!
sub vcl_miss {
    set bereq.http.X-Foo = stale_obj.http.X-Bar;
}

# Always check stale_obj for existence
sub vcl_miss {
    if (stale_obj) {
        set bereq.http.X-Foo = stale_obj.http.X-Bar;
    }
}
```

- No way to prevent this error
 - SHM log records a VCL_error message
- Does stale_obj need to be exposed?

vcl_stale(), vcl_refresh()



```
vcl_stale() ↔ vcl_miss() vcl_refresh() ↔ vcl_fetch()
```

- stale_obj always exists
 - stale_obj could just be obj
- Two new vcl_subs()?

Generalized vcl_stale()



- Either the grace or stale/keep scenario
 - Alternative to vcl_hit() & vcl_miss()
 - Just one timer after ttl?
- Idea: Pull C-logic into vcl
 - pseudocode for HSH_Lookup (after searching for the "best" object)

```
if (oc && (now < ttl))
    next step is HIT;
else if (oc && (now < ttl + keep))
    next step is STALE;
else
    next step is MISS;</pre>
```

vcl_stale() in default.vcl



```
sub vcl_stale {
    /* existing grace logic */
    if (req.busy || !req.backend.healthy)
        return (deliver);
    return (fetch)
}
```

- deliver: deliver the stale object
- fetch: attempt a (possibly conditional) fetch
- pass, error, restart: as usual

vcl_stale()



```
sub vcl_stale {
    /* only part of the current grace logic */
    if (!req.backend.healthy)
        return(deliver);
    /* override conditional requests */
    if (req.url ~ "^/always-fetch") {
        unset bereq.http.If-Modified-Since;
        unset bereq.http.If-None-Match;
        return(fetch);
    /* deliver for restarts */
    if (req.url ~ "^/app" && req.restarts > 0)
        return (deliver);
```

vcl_refresh()?



Expose the stale object after fetch?

If so ...

```
sub vcl_fetch {
    if (stale_obj && stale_obj.hits > 10000)
        set beresp.ttl = 30m;
/* ... or ... ? */
sub vcl_refresh {
    if (obj.hits > 1000)
        set beresp.ttl = 30m;
```



Thanks!

Questions?

- nils.goroll@uplex.de, +49-170-2723133
- geoff.simmons@uplex.de, +49-176-63690917

