



Lesser known VMODs

(some of them)

VUGX

December 3rd 2015

Geoffrey Simmons

Nils Goroll

Agenda

- **Who we are**
- **vmod_re**
- **vmod_vslp**
- **vmod_dcs**
- **others**

Who we are

- 5+ tech people company
 - 3 C-Hackers
- Understand → optimize, fix, secure
- Varnish: Support, Ops, Consulting, Development (vmods, core)
- Simple business model: €/h
- Most outdated website, no twitter, no blogs, no frills

vmmod_re

regex matching and backref capture

```
# Standard backref capture with regsub()
# Write the value of the foo cookie to the Foo header

sub vcl_recv {
    if (req.http.Cookie ~ "\bfoo\s*=\s*\w+") {
        set req.http.Foo
            = regsub(req.http.Cookie, ".*\bfoo\s*=\s*(\w+).*$", "\1");
    }
}
```

- Two regex matches (condition and regsub)
- regsub() must match the entire string
 - '.*' at start may cause deep backtracking
- Cumbersome and verbose

```
import re;

sub vcl_init {
    new foomatcher = re.regex("\bfoo\s*=\s*(\w+)");
}

sub vcl_recv {
    if (foomatcher.match(req.http.Cookie)) {
        set req.http.Foo = foomatcher.backref(1, "fallback");
    }
}
```

backref() : n^{th} subexpression from most recent call to match()

```
import re;

sub vcl_backend_response {
    if (bematcher.match_dyn(beresp.http.URL-Regex, bereq.url)) {
        set beresp.http.Foo = bematcher.backref(1, "fallback");
    }
}
```


match_dyn() compiles the regex in the first parameter on every call

<https://code.uplex.de/uplex-varnish/libvmod-re>


VSLP

- **Main motivation**
 - Sharding
 - Clustering
- **The Problem with modulo hashing**
- **Consistent Hashing**
- **More**

Round-robin



Backend sharding



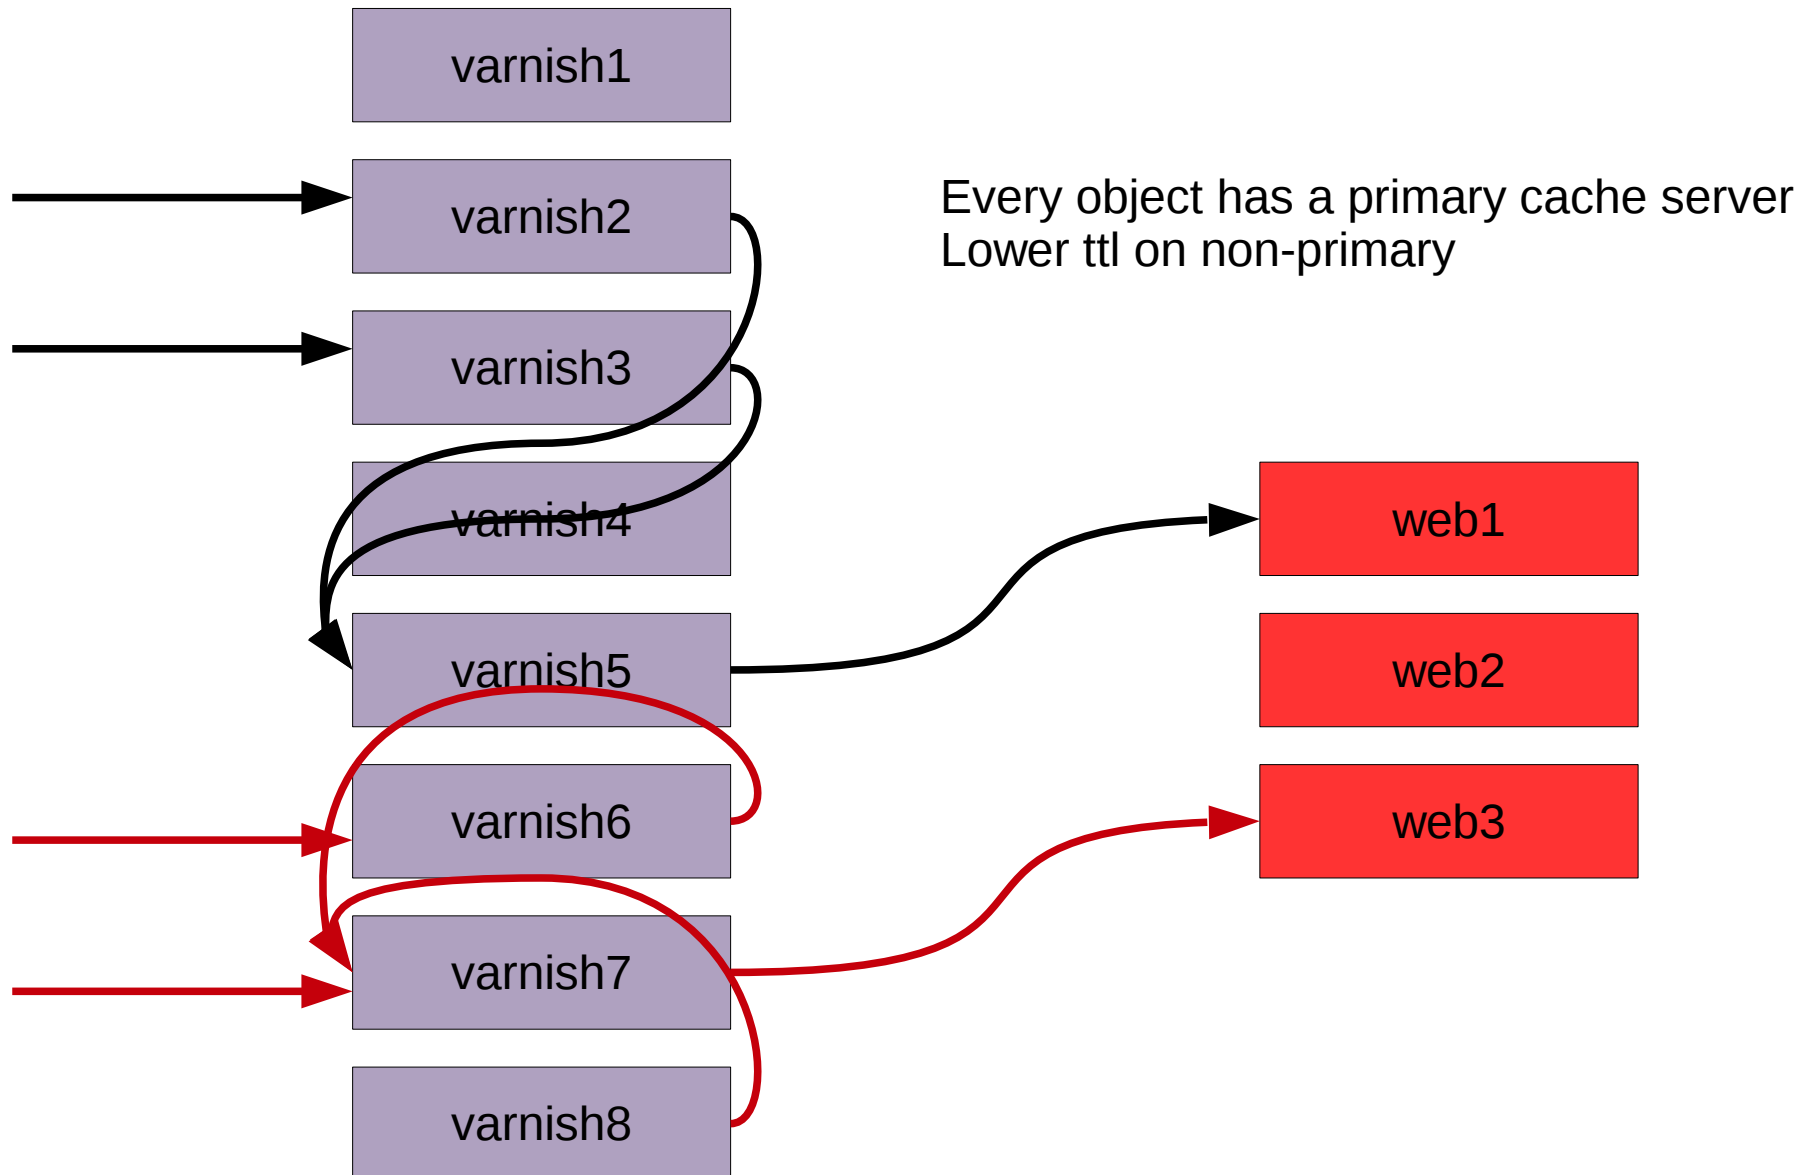
- Hash by URL / ID etc.
- Can be done with hash director

- **Standard hash director:**

`backend = backends[hash % n]`

- **If number/health state of backends changes, most/all hashes get a new backend**
- **Consistent hashing:**
 - **Minimal disturbance of sharding to backends**

Clustering



- **Consistent hashing**
- **Hash with CRC32, SHA256, RS**
- **Rampup**
- **pre-warm aka altsrv_p**
- **Use nth-backend from circle :
restarts**
- <https://code.uplex.de/uplex-varnish/libvmod-vslp>

vmmod_dcs

VMOD dcs

device classifier

```
import dcs;

sub vcl_recv {
    # Get a type name derived from User-Agent
    set req.http.X-UA-Type = dcs.type_name(dcs.classify());

    # Get a meta class for the derived type
    set req.http.X-UA-Class = dcs.type_class(dcs.classify());
}
```

- **maps User-Agent patterns to types and meta-classes**
- **C code is generated from the pattern DB to form a custom matcher.**
 - **scans the string once from left to right**
- **Benchmarks: > 200,000 matches/s (5 μ s/match)**

VMOD dcs

pattern DB



```
# FORMAT: <pattern> <tab> <type>
#
# substr  type
# substr*substr*substr type
# !substr*substr type
```

mozilla*!android*tablet;*firefox/	Tablet
mozilla*android 2.1*movistar link	Mobile Phone

- **Use cases in production:**
 - Device detection
 - Browser version compatibility check
- **Supports Varnish versions 2, 3 and 4**

https://code.uplex.de/uplex-varnish/dcs_classifier

Thanks!

Questions?

- nils.goroll@uplex.de, +49-170-2723133
- geoff.simmons@uplex.de, +49-176-63690917



extra slides


Clustering benefits

- **For n servers**
 - **Effective cache : x times larger**
 - **$1/n$ backend requests**
 - **...**

- **Can be done with the hash director**
- **Cluster:**
 - **Define director of varnish servers**
 - **Select backend**
 - **If self → go to real backend**
 - **Otherwise, recurse**
 - **Lower ttl for recursive requests**

Consistent Hashing

- **Minimizes re-mapping when backends change health / are added/removed**



Actual implementation uses many more nodes

- **More applications:**

- **Hash by IP:**
 - Simple IP persistence
- **Hash by constant value**
 - Priority Server
- **Hash by cache server name (in a cluster)**
 - Each cache server uses a different backend
- ... how many more to come?